

HIDDEN MARKOV MODELS IN SPEECH RECOGNITION

Wayne Ward

**Carnegie Mellon University
Pittsburgh, PA**

Acknowledgements

**Much of this talk is derived from the paper
"An Introduction to Hidden Markov Models",
by Rabiner and Juang**


and from the talk

**"Hidden Markov Models: Continuous Speech
Recognition"
by Kai-Fu Lee**

Topics

- **Markov Models and Hidden Markov Models**
- **HMMs applied to speech recognition**
 - **Training**
 - **Decoding**

Speech Recognition



ML Continuous Speech Recognition

Goal:

Given acoustic data $A = a_1, a_2, \dots, a_k$

Find word sequence $W = w_1, w_2, \dots, w_n$

Such that $P(W | A)$ **is maximized**

Bayes Rule:

$$P(W | A) = \frac{\overset{\text{acoustic model (HMMs)}}{P(A | W)} \cdot \overset{\text{language model}}{P(W)}}{P(A)}$$

$P(A)$ is a constant for a complete sentence


Markov Models

Elements:

States :

$$\mathbf{S} = \{S_0, S_1, \dots, S_N\}$$

Transition probabilities : $P(q_t = S_i \mid q_{t-1} = S_j)$



Markov Assumption:


Transition probability depends only on current state

$$P(q_t = S_i \mid q_{t-1} = S_j, q_{t-2} = S_k, \dots) = P(q_t = S_i \mid q_{t-1} = S_j) = a_{ji}$$

$$a_{ji} \geq 0 \quad \forall j, i$$

$$\sum_{i=0}^N a_{ji} = 1 \quad \forall j$$

Single Fair Coin



$$P(H) = 1.0$$

$$P(T) = 0.0$$

$$P(H) = 0.0$$

$$P(T) = 1.0$$

Outcome head corresponds to state 1, tail to state 2

Observation sequence uniquely defines state sequence

Hidden Markov Models

Elements:

States


$$S = \{S_0, S_1, \dots, S_N\}$$

Transition probabilities


$$P(q_t = S_i \mid q_{t-1} = S_j) = a_{ji}$$

**Output prob distributions
(at state j for symbol k)**

$$P(y_t = O_k \mid q_t = S_j) = b_j(k)$$




Discrete Observation HMM



$$P(R) = 0.31$$

$$P(B) = 0.50$$

$$P(Y) = 0.19$$




$$P(R) = 0.50$$

$$P(B) = 0.25$$

$$P(Y) = 0.25$$

...



$$P(R) = 0.38$$

$$P(B) = 0.12$$

$$P(Y) = 0.50$$

Observation sequence: R B Y Y ... R



not unique to state sequence

HMMs In Speech Recognition

Represent speech as a sequence of observations

Use HMM to model some unit of speech (phone, word)

Concatenate units into larger units



HMM Problems And Solutions

Evaluation:

- Problem - Compute Probability of observation sequence given a model
- Solution - **Forward Algorithm** and **Viterbi Algorithm**

Decoding:

- Problem - Find state sequence which maximizes probability of observation sequence
- Solution - **Viterbi Algorithm**

Training:

- Problem - Adjust model parameters to maximize probability of observed sequences
- Solution - **Forward-Backward Algorithm**

Evaluation

Probability of observation sequence $\mathbf{O} = O_1 O_2 \dots O_T$

given HMM model λ is :

$$P(O | \lambda) = \sum_{\forall Q} P(O, Q | \lambda) \quad Q = q_0 q_1 \dots q_T \text{ is a state sequence}$$

$$= \sum a_{q_0 q_1} b_{q_1}(O_1) \cdot a_{q_1 q_2} b_{q_2}(O_2) \dots a_{q_{T-1} q_T} b_{q_T}(O_T)$$

Not practical since the number of paths is $O(N^T)$

N = number of states in model

T = number of observations in sequence

The Forward Algorithm

$$\alpha_t(j) = P(O_1 O_2 \cdots O_t, q_t = S_j \mid \lambda)$$

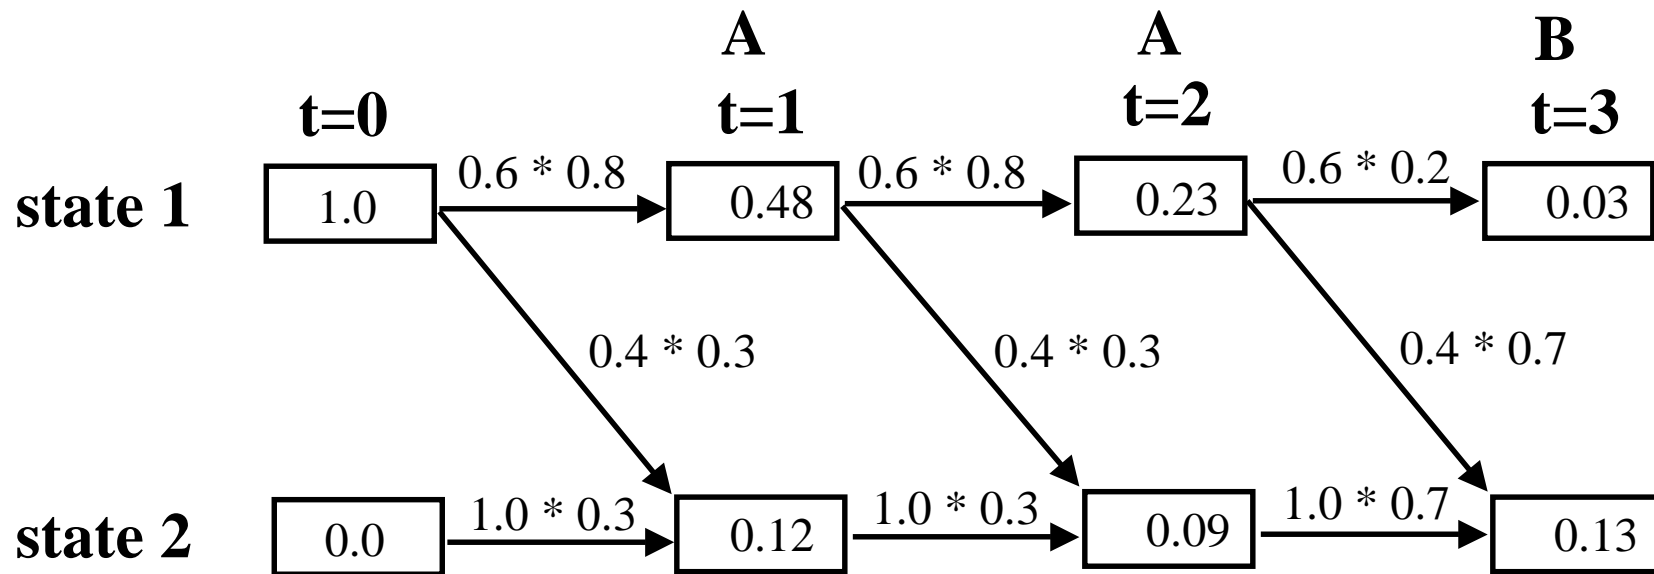

Compute α recursively:

$$\alpha_0(j) = \begin{cases} \mathbf{1} & \text{if } j \text{ is start state} \\ \mathbf{0} & \text{otherwise} \end{cases}$$

$$\alpha_t(j) = \left[\sum_{i=0}^N \alpha_{t-1}(i) a_{ij} \right] b_j(O_t) \quad t > 0$$

$$P(O \mid \lambda) = \alpha_T(S_N) \quad \text{Computation is } O(N^2T)$$

Forward Trellis



The Backward Algorithm

$$\beta_t(i) = P(O_{t+1} O_{t+2} \cdots O_T, q_t = S_i \mid \lambda)$$

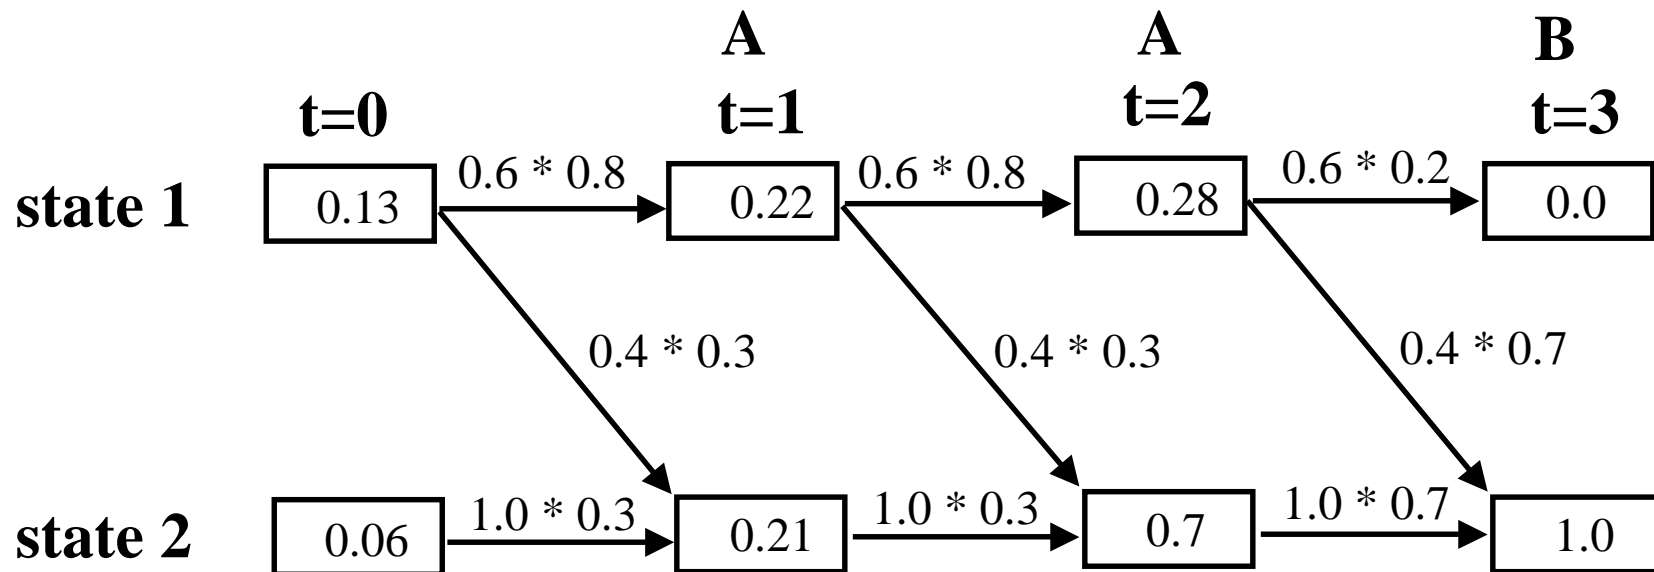

Compute β recursively:

$$\beta_T(i) = \begin{cases} \mathbf{1} & \text{if } i \text{ is end state} \\ \mathbf{0} & \text{otherwise} \end{cases}$$

$$\beta_t(i) = \sum_{j=0}^N a_{ij} b_j(O_{t+1}) \beta_{t+1}(j) \quad t < T$$

$$P(O \mid \lambda) = \beta_0(S_0) = \alpha_T(S_N) \quad \text{Computation is } O(N^2T)$$

Backward Trellis



The Viterbi Algorithm

For decoding:

Find the state sequence **Q** which maximizes $P(\mathbf{O}, \mathbf{Q} \mid \lambda)$

Similar to Forward Algorithm except **MAX** instead of **SUM**

$$VP_t(i) = \text{MAX}_{q_0, \dots, q_{t-1}} P(O_1 O_2 \dots O_t, q_t=i \mid \lambda)$$

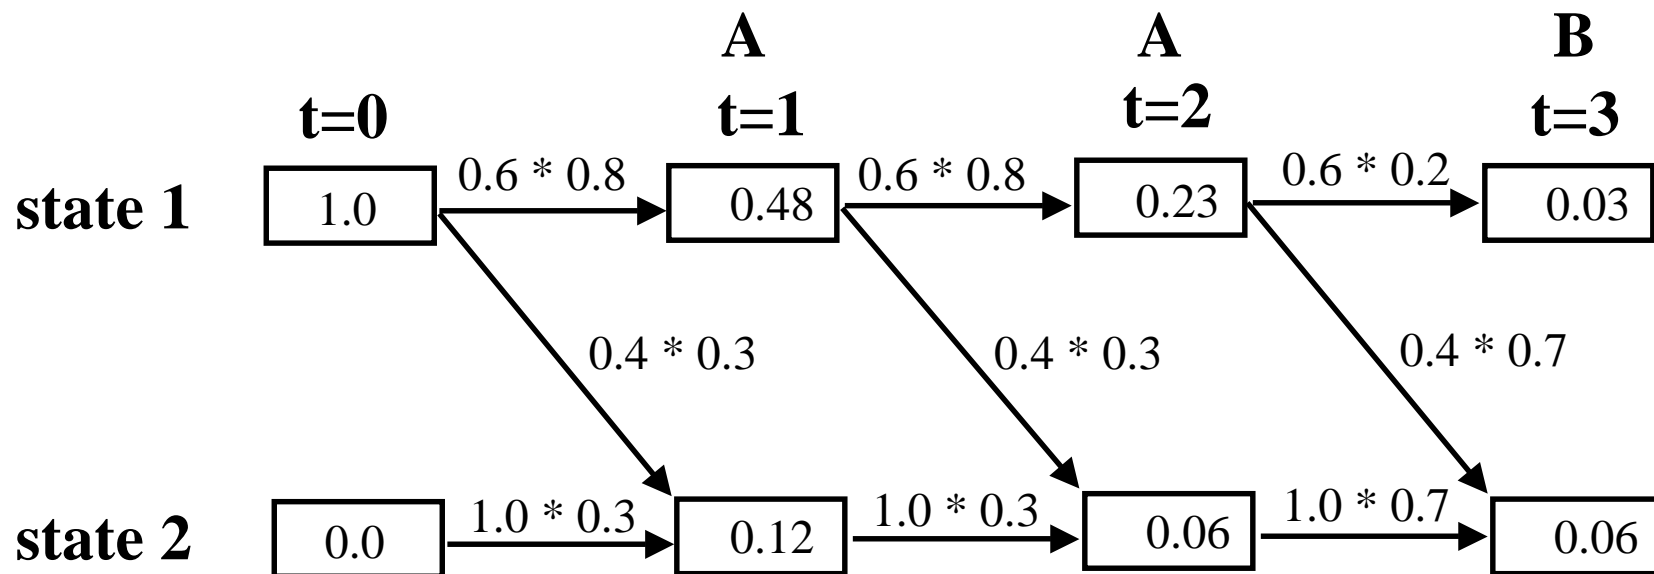

Recursive Computation:

$$VP_t(j) = \text{MAX}_{i=0, \dots, N} VP_{t-1}(i) a_{ij} b_j(O_t) \quad t > 0$$

$$P(\mathbf{O}, \mathbf{Q} \mid \lambda) = VP_T(S_N)$$

Save each maximum for backtrace at end

Viterbi Trellis



Training HMM Parameters

Train parameters of HMM

- Tune λ to maximize $P(O \mid \lambda)$
- No efficient algorithm for global optimum
- Efficient iterative algorithm finds a local optimum

Baum-Welch (Forward-Backward) re-estimation


- Compute probabilities using current model λ
- Refine $\lambda \longrightarrow \hat{\lambda}$ based on computed values
- Use α and β from Forward-Backward

Forward-Backward Algorithm

$\xi_t(i,j)$ = Probability of transiting from S_i to S_j
at time t given O

$$= P(q_t = S_i, q_{t+1} = S_j \mid O, \lambda)$$

$$= \frac{\alpha_t(i) a_{ij} b_j(O_{t+1}) \beta_{t+1}(j)}{P(O \mid \lambda)}$$



Baum-Welch Reestimation

$$\bar{a}_{ij} = \frac{\text{expected number of trans from } S_i \text{ to } S_j}{\text{expected number of trans from } S_i}$$

$$= \frac{\sum_{t=0}^{T-1} \xi_t(i, j)}{\sum_{t=0}^{T-1} \sum_{j=0}^N \xi_t(i, j)}$$

$$\bar{b}_j(k) = \frac{\text{expected number of times in state } j \text{ with symbol } k}{\text{expected number of times in state } j}$$

$$= \frac{\sum_{t: O_t=k} \sum_{i=0}^N \xi_t(i, j)}{\sum_{t=0}^{T-1} \sum_{i=0}^N \xi_t(i, j)}$$

Convergence of FB Algorithm

1. Initialize $\lambda = (A, B)$
2. Compute α , β , and ξ
3. Estimate $\bar{\lambda} = (\bar{A}, \bar{B})$ from ξ
4. Replace λ with $\bar{\lambda}$
5. If not converged go to 2

It can be shown that $P(O | \bar{\lambda}) > P(O | \lambda)$ unless $\bar{\lambda} = \lambda$


HMMs In Speech Recognition

Represent speech as a sequence of symbols

Use HMM to model some unit of speech (phone, word)

Output Probabilities - Prob of observing symbol in a state

Transition Prob - Prob of staying in or skipping state

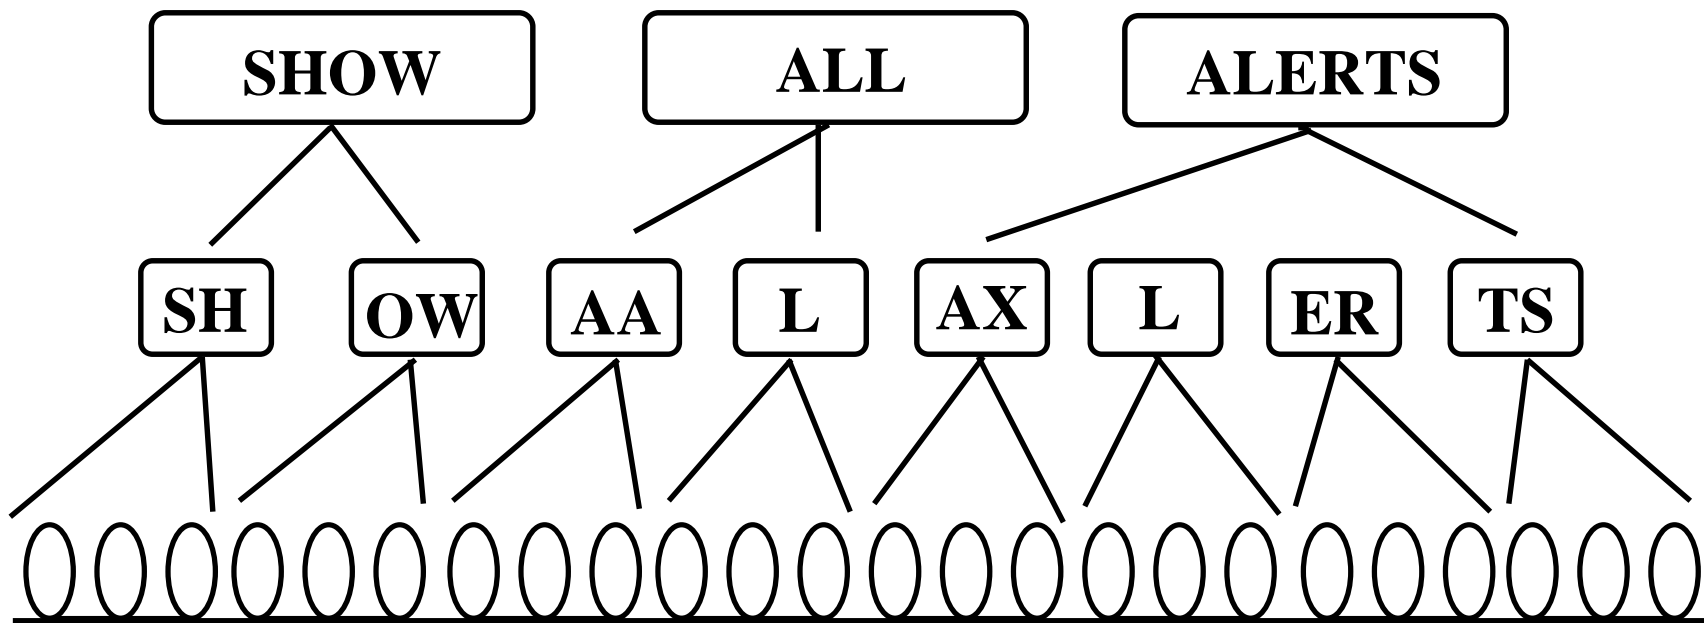


Phone Model

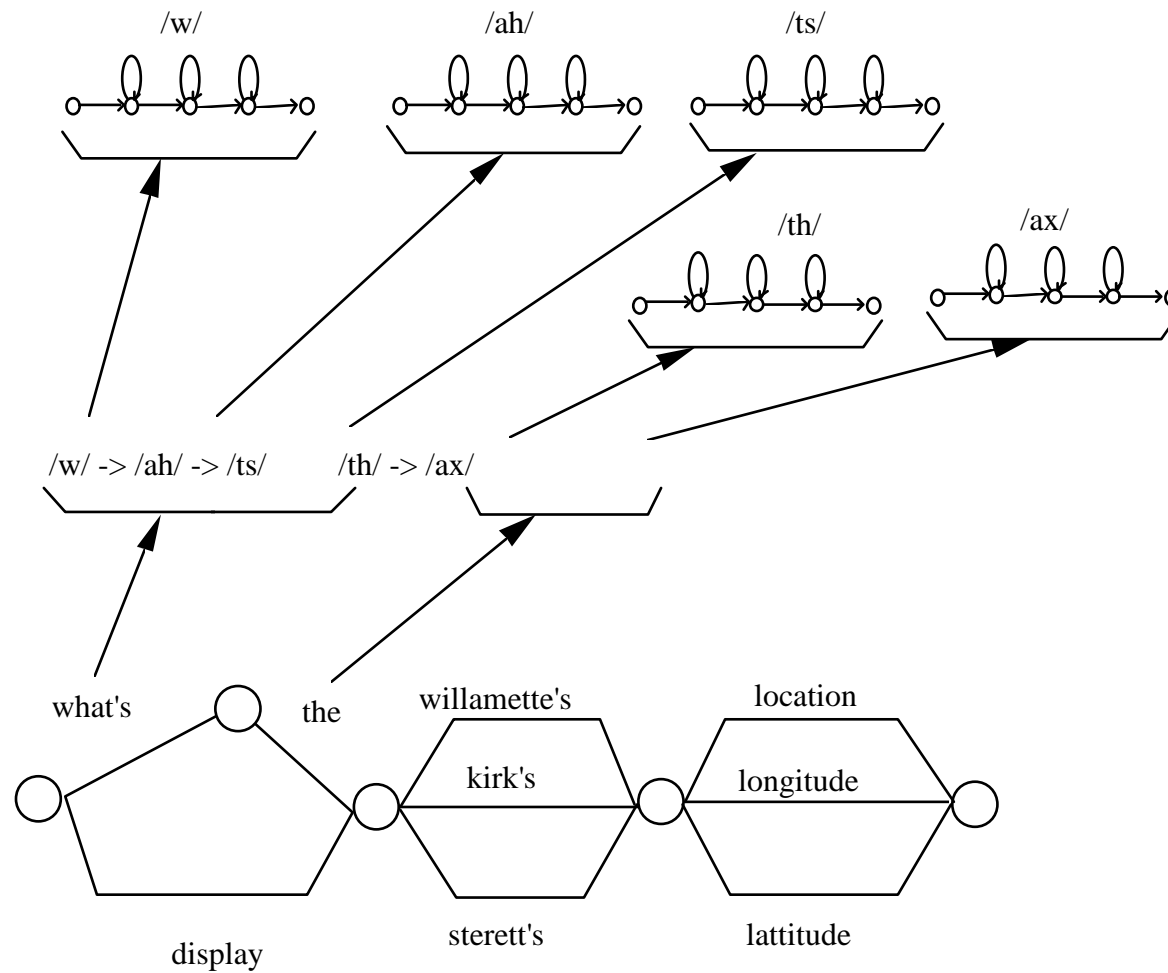
Training HMMs for Continuous Speech

- Use only orthograph transcription of sentence
 - no need for segmented/labelled data
- Concatenate phone models to give word model
- Concatenate word models to give sentence model
- Train entire sentence model on entire spoken sentence

Forward-Backward Training for Continuous Speech



Recognition Search




Viterbi Search

- **Uses Viterbi decoding**
 - **Takes MAX, not SUM**
 - **Finds optimal state sequence $P(O, Q | \lambda)$
not optimal word sequence $P(O | \lambda)$**
- **Time synchronous**
 - **Extends all paths by 1 time step**
 - **All paths have same length (no need to
normalize to compare scores)**

Viterbi Search Algorithm


0. Create state list with one cell for each state in system
1. Initialize state list with initial states for time $t = 0$
2. Clear state list for time $t+1$
3. Compute within-word transitions from time t to $t+1$
 - If new state reached, update score and BackPtr
 - If better score for state, update score and BackPtr
4. Compute between word transitions at time $t+1$
 - If new state reached, update score and BackPtr
 - If better score for state, update score and BackPtr
5. If end of utterance, print backtrace and quit
6. Else increment t and go to step 2

Viterbi Search Algorithm



$\text{OldProb}(S1) \cdot \text{OutProb} \cdot \text{Transprob}$

$\text{OldProb}(S3) \cdot P(W2 \mid W1)$



Viterbi Beam Search

Viterbi Search

All states enumerated

Not practical for large grammars

Most states inactive at any given time

Viterbi Beam Search - prune less likely paths

States worse than threshold range from best are pruned

From and To structures created dynamically - list of active states

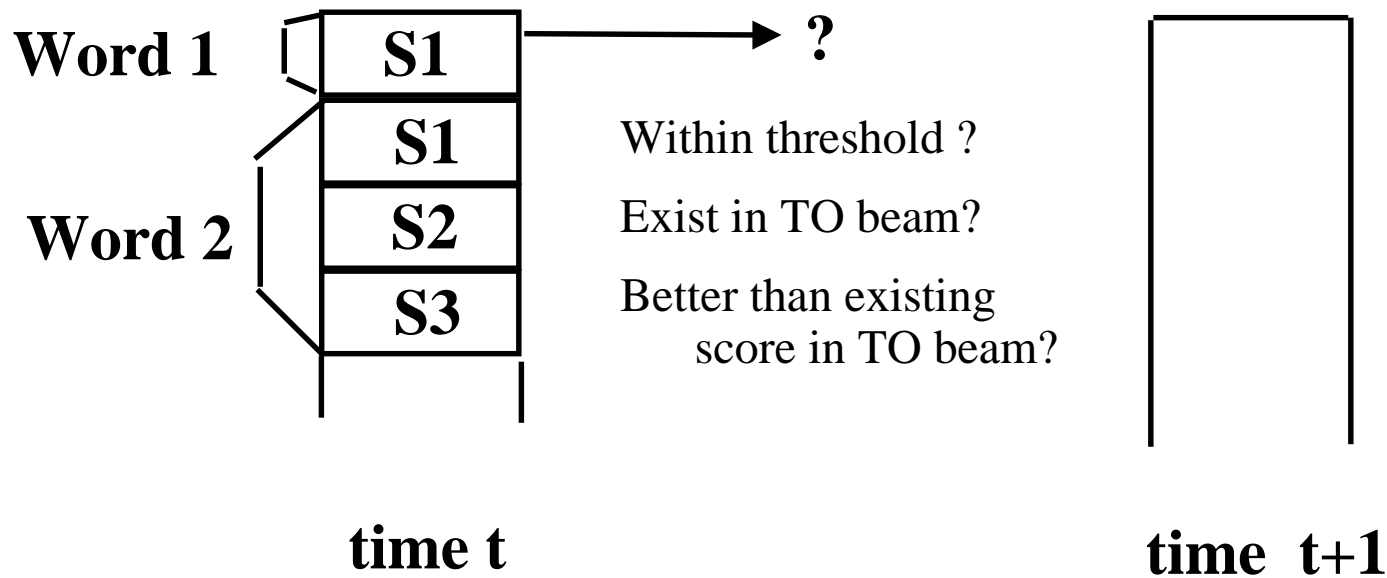
Viterbi Beam Search

**FROM
BEAM**

States within threshold
from best state

**TO
BEAM**

Dynamically constructed



Continuous Density HMMs

Model so far has assumed discrete observations,
each observation in a sequence was one of a set of M
discrete symbols

Speech input must be Vector Quantized in order to
provide discrete input.

VQ leads to quantization error

The discrete probability density $b_j(k)$ can be replaced
with the continuous probability density $b_j(\mathbf{x})$
where \mathbf{x} is the observation vector

Typically Gaussian densities are used

A single Gaussian is not adequate, so a weighted sum of
Gaussians is used to approximate actual PDF

Mixture Density Functions

$b_j(\mathbf{x})$ is the probability density function for state j

$$b_j(x) = \sum_{m=1}^M c_{jm} N[x, \mu_{jm}, U_{jm}]$$

\mathbf{x} = Observation vector x_1, x_2, \dots, x_D

M = Number of mixtures (Gaussians)

c_{jm} = Weight of mixture m in state j where $\sum_{m=1}^M c_{jm} = 1$

N = Gaussian density function

μ_{jm} = Mean vector for mixture m , state j

U_{jm} = Covariance matrix for mixture m , state j

Discrete Hmm vs. Continuous HMM

□ Problems with Discrete:

- quantization errors
- Codebook and HMMs modelled separately

• Problems with Continuous Mixtures:


- Small number of mixtures performs poorly
- Large number of mixtures increases computation and parameters to be estimated

$$c_{jm}, \mu_{jm}, U_{jm} \text{ for } j = 1, \dots, N \text{ and } m = 1, \dots, M$$

- Continuous makes more assumptions than Discrete, especially if diagonal covariance pdf
- Discrete probability is a table lookup, continuous mixtures require many multiplications

Model Topologies

Ergodic - Fully connected, each state has transition to every other state



Left-to-Right - Transitions only to states with higher index than current state. Inherently impose temporal order. These most often used for speech.

